

Your Index Is Fine, Your Query Encoder Is Not: Query Distillation over Frozen Indexes

Dylan Couzon	Evgeniya Sukhodolskaya
dylan.couzon@qdrant.com	jenny@qdrant.com
Kumar Shivendu	Andrey Vasnetsov
kumar.shivendu@qdrant.com	andrey@qdrant.com
Qdrant	

October 2026

Abstract

The cost of a large document encoder can be amortized across many searches using the same stored vectors. Query encoding adds compute to each uncached query, even if that query is used only once. Changing the document encoder requires re-embedding the corpus and rebuilding its index. Query-side distillation offers a cheaper query encoder, or student, that searches the existing document vectors. What retrieval quality can that student preserve, and how much of its encoding saving survives approximate search?

We fit two inexpensive query students across 26 embedding spaces under one protocol: a token-vector table that sums its tokens’ vectors, and a linear projection of a small frozen transformer’s features. At a given vector width, stronger original encoders tend to yield stronger students; after accounting for encoder quality, wider vectors predict weaker students. A model using both properties predicts student rankings on held-out and prospectively tested encoders better than encoder quality alone. The prospective evidence is stronger for the transformer-based student.

The token-table students require a median two to four times the graph-search budget across three corpora to match the original query path’s fraction of recovered exact neighbors, with each path measured against its own exact results. The corresponding relevance-loss penalty depends on the corpus. The neighbor-recovery penalty also appears in LightRetriever’s jointly trained lookup path. Constella demonstrates the substitution over Stella, a frozen document encoder: its compact transformer and token-table query paths retain 90.5% and 81.4% of Stella’s exact retrieval quality on BEIR-15. Separate timing measurements show that extra search narrows their encoding advantage. Student quality and encoding-plus-search cost together determine the value of a cheaper query path over an existing index.

1. Introduction

In a production retrieval system, reducing query compute by replacing the embedding model can turn an optimization into an index migration. Changing the document representation requires re-embedding the corpus and rebuilding the vector index. Downtime can be avoided, but keeping retrieval available during the transition can require old and new vector indexes to coexist and ingestion to keep both representations current until traffic switches [Index migration]. Backfill and index construction then add compute, storage, and I/O demands alongside live queries and updates. The migration also needs relevance validation, a coordinated query-encoder switch, and a rollback path. Its duration and cost depend on corpus token volume, encoding throughput, index-building capacity, and the rate of incoming updates. Replacing only the query encoder with a compatible student preserves the stored document vectors and avoids that migration; its effects on retrieval quality and search effort still need to be measured.

Query-side distillation trains a student encoder to reproduce the frozen encoder’s query vectors and search the frozen document vectors directly [QED, EmbedDistill, LEAF]. Static variants fit a token table, so a query costs a table lookup and a sum [pyNIFE, LightRetriever]. We study the retrieval quality and search effort of these replacements, and whether properties of the existing embedding space predict their performance.

We examine two properties of the space: vector width and the original encoder’s retrieval quality. A frozen index inherits both from its document encoder. Fitting each student by a linear solve takes minutes, allowing the same recipe to be compared across 26 public encoder spaces. Two questions organize the paper.

- **RQ1, retention.** How much retrieval quality can a query student fitted by a linear solve recover from a frozen index, and do vector width and the space’s own quality predict it? We test the project’s initial heuristic that a better space yields a better student, using exact retrieval.
- **RQ2, search effort.** Over the index’s unchanged approximate-search graph, how much more search effort does the student’s query path need than the index’s own query path, and is that effort predictable?

Student quality follows two opposing associations: stronger spaces favor the student, while wider vectors work against it. A model using both properties predicts held-out and prospectively tested encoders (Section 4). The token-table students also need more graph effort to recover their exact neighbors in every measured space on two of three corpora. The distance between queries and documents predicts the size of that penalty, which also appears in LightRetriever’s jointly trained lookup path (Section 5).

Constella is the worked instance: a frozen Stella index [Stella] with three compatible query paths, Stella’s own encoder, a 34.5-million-parameter transformer student (Nano), and a token-table student (Zero), released with the registered evaluations that gated the release. Section 6 reports it.

Throughout, “registered” marks the project’s pre-registered evaluations, whose decision rules were fixed before any result was observed, and “exploratory” marks analyses added afterwards. The two index-wide screens that answer RQ1 and RQ2 are exploratory; their analyses were written down before scoring, and the prospective test of RQ1 commits its predictions before any student is fitted.

2. Related work

Query-side distillation against a frozen document encoder is established: QED retains 92.5% of a dense retriever’s BEIR score with a two-layer student [QED], EmbedDistill distills an asymmetric student against the teacher’s frozen document encoder [EmbedDistill], and LEAF reports 97.7% retention at 4.7-fold compression [LEAF]. pyNIFE fits a static token table to a frozen teacher [pyNIFE], while LightRetriever trains a lookup query side jointly with a document tower [LightRetriever]. Static sentence embeddings, Model2Vec, and NanoVDR explore related compact representations [StaticEmb, Model2Vec, NanoVDR]. RQ1 extends this work with a common-recipe comparison across encoder spaces; RQ2 tests LightRetriever’s released lookup path over its own index alongside the fitted students.

Stronger teachers can produce weaker students in distillation generally [Cho and Hariharan 2019] and in dense retrieval [PROD]; those works discuss capacity gaps between teacher and student. Here, vector width is a predictive property tested across encoder spaces, with held-out and prospective predictions.

Graph-based approximate search over HNSW is sensitive to where queries sit relative to the indexed vectors [HNSW]. Out-of-distribution queries can need more effort, and graph constructions use query samples to mitigate it [OOD-DiskANN, RoarGraph]. RQ2 examines query substitution over an unchanged graph. Query-aware decoding of binary sketches is the precedent for the precision control in Appendix C [QA-Cos].

3. Experimental setup

3.1 Encoder spaces

The roster is 26 public English embedding encoders that share the 30,522-entry BERT WordPiece vocabulary, a requirement of the table student. Ten were registered before any result; 16 were added as an exploratory

expansion under the same recipe and selection order. Each encoder defines one frozen space: its document vectors are frozen and its own query encoder is the reference path. Table A lists them by width.

Table A. The 26 encoder spaces by width (full roster in Appendix B). “Own nDCG@10” is the encoder’s own query path on the six-set suite (Section 3.3); student columns are the outcomes of Section 4.

Width	Spaces (registered)	Own nDCG@10, range	Table student, range	Head student, range
384	9 (0)	0.328 to 0.504	0.290 to 0.358	0.196 to 0.508
768	10 (4)	0.285 to 0.533	0.186 to 0.353	0.059 to 0.325
1024	7 (6)	0.474 to 0.597	0.246 to 0.397	0.156 to 0.268

The roster spans nine related model families and three width levels; several checkpoints are revisions of one another. We report checkpoint-resampling intervals alongside family-level sensitivity analyses. The arctic-embed-l mean-pooling control is excluded from roster statistics; its published CLS pooling is used for the roster entry.

3.2 Student representations

Two closed-form students, fitted by linear solves, produce query vectors at each index’s width.

Table student. Ridge regression fits one vector per WordPiece token to reproduce the original encoder’s query vectors. At query time, token vectors are summed with square-root token-count weights and the result is normalized. Encoding costs a tokenizer pass and a vector sum. Initialization and solver details are in Appendix B.

Head student. The small bge-small transformer provides frozen query features: the masked-mean states of layers 12, 8, and 4, concatenated into 1,152 dimensions plus a bias. A ridge head maps these features into the index’s space and normalizes the output. This second representation tests whether the findings extend beyond the token table; Nano uses the same feature path.

These fits provide inexpensive baselines for the two representations. Section 6 evaluates the separately trained Zero and Nano models.

3.3 Fit data, selection, and scoring

Both students are fitted on the same 337,981 screened fit queries, encoded by each index’s own query path with its published prompt. The ridge penalty is chosen per index on two development forums (CQADupStack physics and programmers) and frozen before any public dataset is scored. Retrieval is scored by exact search, nDCG@10 with the registered self-hit drop, on six public BEIR datasets [BEIR]: NFCorpus, SCIDOCS, SciFact, TREC-COVID, FiQA, and ArguAna. The six-set macro is the outcome; the four sets without disclosed exposure for the Stella instance (the first four) are a sensitivity outcome. Each index is scored once per set.

3.4 Approximate-search protocol

RQ2 uses Qdrant 1.19.1 [Qdrant 1.19.1] with uncompressed vectors, cosine distance, and HNSW graphs (m=16, ef_construct=100). We build two graphs per space and workload using different seeded insertion orders. The workloads are FiQA (57,638 documents), SCIDOCS (25,657), and TREC-COVID (171,332). Each graph is searched with the original encoder’s queries and both students’ queries, sweeping ef, which controls the breadth of graph search, from 16 to 512.

Each path is compared with its own exact top-10. The pre-specified primary measure is the relative-nDCG-loss multiplier: the smallest student ef that reaches the original path’s relative loss at ef=64, divided by 64 and averaged over builds. A student that fails to reach that reference by ef=512 in a build is reported as censored, with that build’s multiplier at least eight. The secondary measures are the analogous multiplier for

tie-aware recovery@10, the fraction of exact neighbors recovered, and the recovery gap at `ef=64`. Appendix C gives the parity checks, tie rule, and tas-b exclusion.

3.5 The instance

Constella is built around Stella, the `stella_en_400M_v5` encoder, whose 1024-dimensional normalized document vectors are frozen throughout; Zero is a trained table student and Nano a 34.5-million-parameter transformer on the head student’s feature path, both trained against Stella’s query vectors and evaluated on BEIR-15 by exact search (Section 6, Appendix A).

4. RQ1: what the index decides about retention

4.1 Hypothesis and the pooled result

H1. The stronger the index’s own retrieval, the stronger the closed-form student fitted to it; in particular, the strongest index yields the strongest student.

H1 was the project’s initial selection rule. The strongest registered space, `gte-large-en-v1.5`, yields the weakest registered table student and a below-median head student. Stella, second in the same width band, yields the strongest table student. Across all 26 spaces, the rank correlation between the original encoder’s score and its student’s score is +0.09 for the table (95% checkpoint interval -0.39 to $+0.52$) and +0.09 for the head (-0.36 to $+0.51$). These wide intervals allow a moderate positive association. Section 4.4 evaluates the selection rule by the quality of the student it picks.




Figure 1: Student quality against the space’s own quality for the table (left) and head (right) students, colored by width. Filled points are registered spaces; triangles are the prospective spaces of E24 (eight for the head, seven for the table).

4.2 Width as the competing explanation

H1’. At a given width, stronger spaces yield stronger students; after accounting for space quality, wider vectors predict weaker students. Support requires a width coefficient whose interval excludes zero and better held-out-family prediction than space quality alone.

Table B models the student’s absolute six-set score using the space’s own score and $\log_2$ vector width. Retention is the ratio of student to original-encoder scores. Standardized coefficients express the change in

student quality, in standard deviations, per standard deviation of a predictor. Intervals use 10,000 checkpoint resamples. Leave-one-family-out prediction fits on eight model families and predicts the ninth.

Table B. Nested models of student quality over 26 indexes (exploratory, analysis pre-specified).

Student	Predictors	Space quality beta	Width beta	In-sample R^2	Leave-one-family- out Spearman
Head	space quality	+0.34 [−0.16, +0.65]		0.12	−0.17
Head	space quality, width	+0.69 [+0.40, +0.97]	−0.83 [−1.14, −0.60]	0.70	+0.69
Table	space quality	+0.37 [−0.30, +0.72]		0.14	−0.15
Table	space quality, width	+0.63 [+0.06, +0.89]	−0.64 [−0.95, −0.32]	0.48	+0.40

Space quality and width correlate at +0.61 across the roster. Their opposing associations with student quality help explain the near-zero pooled rank correlation. Given width, the partial rank correlation between student and index quality is +0.58 for the head and +0.45 for the table. Table B shows a larger standardized coefficient for width under the head representation and coefficients of similar magnitude under the table representation. Family-adjusted fits and coefficient intervals are reported in Appendix B.

4.3 Held-out and prospective prediction

We test whether the quality-and-width model ranks further encoders better than their own retrieval scores do. Fitted on the ten registered spaces, the model ranks the 16 later spaces at Spearman 0.85 [0.53, 0.99] for the head and 0.70 [0.20, 0.96] for the table; the original scores alone rank them at 0.37 and 0.16. The hypothesis was formed after all 26 spaces had been seen, making this a held-out fit.

For the prospective test, we chose eight previously unscored BERT-vocabulary encoders spanning the three widths, fitted the model on the 26, and committed its predictions by hash before fitting any student (Appendix B). One table fit stopped at the convergence gate, leaving seven table students and eight head students. The committed predictions rank them at Spearman +0.86 [+0.33, +1.00] for the head and +0.71 [−0.12, +1.00] for the table, against +0.26 and +0.39 for the original scores alone. Mean absolute error of the predicted six-set score is 0.039 (head) and 0.032 (table). The head’s checkpoint interval supports a positive association; the table’s wider interval leaves its prospective result uncertain.

4.4 Corollary: choosing an index before it exists

RQ1 concerns a space that already exists. For a team that can still choose, the same data answer a second question: which rule picks the encoder whose student will be best? Table C scores three rules by the six-set nDCG@10 their pick loses against the best student in the roster.

Table C. Six-set nDCG@10 lost against the best student in the roster, by selection rule.

Rule	Head: pick, regret	Table: pick, regret
Strongest space (H1)	gte-large-en-v1.5, 0.285	gte-large-en-v1.5, 0.152
Two-variable model, leave-one-out pick	arctic-embed-s, 0.249	gte-small, 0.107
Student’s own dev-forum screen	gte-small, 0.035	stella-400M-v5, 0.000
Prospective spaces (eight head, seven table): strongest space	0.086	0.107
Prospective spaces: committed model pick	0.000	0.043
Prospective spaces: dev screen	0.000	0.043

The student’s development screen makes the best or near-best pick under both representations. It scores the fitted student on two forums and predicts six-set rank at Spearman 0.84 (head) and 0.88 (table). Adding it to the quality-and-width model gives it the largest standardized coefficient (+0.71 [+0.36, +0.99] for the head) and raises leave-one-family-out prediction to 0.77 and 0.80. The quality-and-width model captures the ranking structure, while testing a fitted student improves selection. On the prospective spaces, the model’s committed pick matches the development screen’s; selecting by the original encoder’s score incurs the largest regret on both rosters.

The head’s backbone, bge-small-en-v1.5, is also a roster encoder. Removing it changes every statistic above by less than 0.06.

5. RQ2: what the index decides about search effort

5.1 Hypothesis

H2. Over the index’s unchanged HNSW graph, the student’s query path needs more search effort than the index’s own query path to recover the same fraction of its exact neighbors, because the student’s queries sit differently against the indexed vectors.

We first measure the effort penalty per index, then test whether query-placement features predict its size (Section 5.3).

5.2 Effort across 25 indexes

H2 predicts a multiplier above one in most spaces on every workload. Table D tests it.

Table D. Search effort of the table student relative to the index’s own path at **ef=64**, two builds averaged, 25 indexes (exploratory, analysis pre-specified). The primary quantity is the relative-nDCG-loss multiplier; “censored” means at least one build never reached the reference by **ef=512**.

Workload	Documents	Loss multiplier: indexes above 1 / at or below	Censored	Loss multiplier median over reached	Recovery gap at ef=64, mean (indexes below 0)	Recovery multiplier median
FiQA	57,638	25 / 0	4	4.0	−2.9 pp (25 of 25)	4.0
SCIDOCS	25,657	17 / 8	8	1.25	−0.8 pp (17 of 25)	2.0
TREC-COVID	171,332	12 / 13	8	0.25	−3.9 pp (25 of 25)	4.0

The primary measure holds on FiQA (25 of 25, median four over the 21 reached, four censored in at least one build) and on SCIDOCS (17 of 25). It is mixed on TREC-COVID (12 of 25): over 50 queries, the index’s own path loses only 0.2% of nDCG at **ef=64**, so the relative-loss threshold there is sensitive to differences smaller than its noise. By the secondary measure it holds on every workload (Table D), and the head student shows the same recovery deficit in 24 of 25 spaces on FiQA and 23 of 25 on TREC-COVID.

H2c predicts a smaller recovery gap on a smaller corpus. On FiQA subsampled to SCIDOCS’s 25,657 documents, with the same queries and two rebuilt graphs per space (exploratory, pre-specified; Appendix C), the gap at **ef=64** moves from −2.9 points to −1.2, a paired change of +1.8 points (95% interval over spaces +1.3 to +2.3). It remains 0.4 points larger than SCIDOCS’s (−0.9 to −0.1). The subsample reproduces much of SCIDOCS’s attenuation, with a remaining difference between the corpora.

5.3 Is the effort predictable?

H2b predicts that query-placement features computed before graph construction rank the spaces by their recovery gap. We declared nine features and two covariates before computing them (exploratory, pre-specified; Appendix C). Table F focuses on the query-to-document distance ratio: the distance from queries to their nearest documents, relative to the typical distance between neighboring documents. Correlations use 75 space-by-workload points, with intervals resampling each space together with its three workload rows. Held-out prediction fits a ridge model on all families except one and predicts that family.




Figure 2: Recovery gap at **ef=64** by index, three workloads, two builds averaged.

Table F. Predictors of the recovery gap over 75 points (25 spaces by three workloads). Negative Spearman means a larger feature value goes with a larger penalty.

Feature	Spearman, 75 points	95% interval	FiQA / SCIDOCS / TREC-COVID	Leave-one-family- out Spearman, MAE (pp)
Query-to- document distance ratio, encoder’s queries	−0.64	[−0.75, −0.48]	−0.80 / −0.32 / −0.83	+0.64, 1.26
Query-to- document distance ratio, student’s queries	−0.63	[−0.74, −0.46]	−0.77 / −0.45 / −0.69	+0.56, 1.79
All nine features and two covariates				+0.73, 1.40
Training-fold mean (baseline)				2.22

Spaces with queries farther from documents tend to pay a larger recovery penalty. The ratio computed from the original encoder’s queries predicts held-out families with lower absolute error than the all-feature model (Table F), so the existing query path already supplies a useful predictor. The association is clearest on FiQA and TREC-COVID; on SCIDOCS its interval includes zero. Appendix C reports the remaining features, including changes between student and original-encoder queries.

Prediction across workloads is less precise. A model fitted on two workloads ranks the third at +0.50, but its mean error is 2.55 points against 2.33 for the training-fold mean. The ratio therefore supports relative ranking, with the magnitude of the penalty requiring calibration on the target workload.

5.4 The penalty in a jointly trained model

H2d. If the deficit comes from where lookup queries sit against a graph built for the encoder’s own queries, it should also appear when the lookup side is trained jointly with the document tower rather than fitted afterwards. We tested the released LightRetriever model (Qwen2.5-1.5B with its published adapter) over its own FiQA and SCIDOCS document vectors, comparing its lookup query path with the same model’s full query encoding under the Section 3.4 protocol, with the model’s web-search instruction as the primary prompt and the task prompt as secondary. Predictions and quantities were fixed before encoding (exploratory, pre-specified).

Table E. LightRetriever’s lookup path against its full query path over its own index, two builds averaged.

Workload	Prompt	Exact nDCG@10, lookup / full	Loss multiplier (builds)	Recovery multiplier	Recovery gap at ef=64, query 95% interval
FiQA	web search	0.4065 / 0.4752	3.0 (2, 4)	4.0	−3.6 pp [−4.3, −2.9]
FiQA	task	0.4124 / 0.4766	2.0 (2, 2)	4.0	−2.3 pp [−2.9, −1.8]
SCIDOCS	web search	0.1655 / 0.1925	2.25 (0.5, 4)	4.0	−1.7 pp [−2.1, −1.4]
SCIDOCS	task	0.1818 / 0.2017	censored (1, censored)	2.0	−0.7 pp [−1.0, −0.5]

The lookup path recovers fewer of its own exact neighbors at **ef=64** on both corpora under both prompts, with query intervals clear of zero. Matching the full path’s recovery takes four times the **ef** in every FiQA build and under the primary SCIDOCS prompt. The relevance-loss multiplier agrees on FiQA (two to four); on SCIDOCS, both paths lose under 0.1% of nDCG at **ef=64**, making that reference sensitive to very small differences. GPU encoding took 0.10 ms for the lookup path and 0.90 ms for the full path. These encoding timings and graph-effort multipliers measure separate components; Section 6 combines encoding and search timing for Constella.

The **ef** multipliers measure graph-search effort. Relevance loss and exact-neighbor recovery give complementary views of the substitution, with different conclusions on TREC-COVID and on LightRetriever’s SCIDOCS index.

6. The Constella instance

Table 1 gives the three query paths over Stella’s frozen index, by exact search over BEIR-15.

Table 1. Three query paths over one frozen Stella index. Encoding medians: 100 real 5-to-12-word queries, batch one, four CPU threads, Apple M5 Pro, registered protocol.

Query path	Exact nDCG@10	Share of Stella	Encode p50	Assets
Stella query encoder	0.5614	reference	31.6 ms	1669.6 MiB
Nano (34.5M transformer)	0.5081	90.5%	2.25 ms	132.3 MiB
Zero (token table)	0.4572	81.4%	0.044 ms	90.1 MiB

Zero sums learned token vectors; Nano uses contextual features from three layers of a small transformer. Both reduce encoding time while preserving a different share of Stella’s exact retrieval quality (Table 1). All three query paths were verified against the same populated collections, allowing each request to choose a compatible encoder. Appendix A gives the per-dataset results and selected-reference comparisons.

Build cost and query cost. Table-student fitting took four to seven minutes of A100 time per encoder, including query encoding except for Stella’s cached targets. Nano’s final optimization took 57.3 A100 hours, about \$95 at the recorded rental rate, excluding data preparation, recipe search, and failed runs. Appendix B gives the component costs. Both constructions preserve the existing document index.

Table 2 extrapolates registered timing medians to encoding-plus-search time per million queries on the measurement laptop. Encoding uses batch one with four threads; search uses the uncompressed million-passage collection at each path’s setting within 1% of its own exact score. These are sums of separately measured encoding and search times.

Table 2. Extrapolated encoding-plus-search time per million queries on the measurement laptop, from the medians in Table 1 and Appendix C.

Query path	Encoding	Search, within 1% of own exact	Total	Share of Stella
Stella query encoder	8.78 h	0.43 h (ef =128)	9.2 h	reference
Nano	0.63 h	0.61 h (ef =256)	1.2 h	13%
Zero	0.01 h	0.99 h (ef =512)	1.0 h	11%

Search time at a fixed **ef** differs by less than 5% across paths. The extra work comes from the larger **ef** needed to reach each path’s relative-quality target: 128 for Stella, 256 for Nano, and 512 for Zero. Zero’s extra search largely consumes its encoding advantage over Nano, bringing their extrapolated encoding-plus-search times close together (Table 2). Their exact BEIR-15 quality is reported separately in Table 1. On one binary-quantized collection, the 63-fold encoding gap between Zero and Nano becomes 2.2-fold after search, at different quality levels (Figure 3). Appendices C and D report query-precision controls, lexical fusion, and per-query routing over the same index.




Figure 3: Quality under approximate search against encoding plus search time, registered sweep. Dotted lines are each path’s exact score.

7. Limitations

RQ1 covers two closed-form recipes over 26 related spaces at three width levels. Width covaries with family and quality, making its association predictive rather than causal. Checkpoint-resampling intervals describe variation within this roster; family-level prediction tests transfer across its families. Transfer to fully trained students such as Zero and Nano remains an open question.

RQ2 covers one engine, one graph configuration, and three workloads with different sizes and domains, including 50 queries on TREC-COVID. Its predictors establish associations with search effort; cross-space timing ran on a shared machine, so latency conclusions rely on the separate Constella sweep. The LightRetriever replication covers one released 1.5-billion-parameter model, our reproduction of its dense path, and two corpora.

Constella’s students use different training pools, including different FEVER exposure, and Stella discloses exposure on ArguAna, FiQA, and FEVER. Query intervals condition on the trained models. Three bounded

specialization attempts on an internal workload produced no improved encoder; per-workload students remain a direction for future research.

Artifact availability

The companion repository holds inference examples, build configurations, methods with their dated amendments, immutable result receipts, per-dataset tables, and the full research history including negative results. The released artifacts are Constella Zero, Constella Nano, and the Stella document encoder on the Hugging Face Hub. Two inputs are identified by hash but not yet archived in the repository: the cleaned 337,981-query fit list and the sampled one-million-passage identifier list.

Appendix A. Registered evaluations of the instance

Serving configuration. Stella is pinned to revision `ffeb2b7e`, with its `s2p_query` prompt, fp32 encoding, normalized output, fp16 stored document vectors, and cosine similarity; maximum sequence length 512. Nano uses bge-small layers 12, 8, and 4 concatenated to 1152 dimensions before projection; batches mix 75% queries and 25% documents. Zero’s training pool includes FEVER-train and HotpotQA; Nano excludes FEVER. MS MARCO is licensed for non-commercial use, so it was validation-only in our distillation: it supplied no training targets, negatives, or generation seeds, although Nano’s pretrained base has prior exposure to it.




Figure A1: Retention per dataset. Datasets with disclosed Stella training exposure are ArguAna, FiQA, and FEVER; the four originally held-out sets are FEVER, DBpedia-entity, and the CQADupStack android and english forums.

Variation across datasets. Zero’s sum of token vectors is insensitive to word order and context. Its retention is 0.67 on TREC-COVID and FiQA, against 0.95 on Quora. Nano’s retention ranges from 0.85 to 0.99 except FEVER, which was in Zero’s training pool and excluded from Nano’s. These dataset differences describe the released models and their distinct training pools.

Partitions. The six-set suite is NFCorpus, SCIDOCs, SciFact, TREC-COVID, FiQA, and ArguAna. “Clean-four” means the first four, which have no disclosed Stella overlap; it is not a causal contamination estimate. Four datasets were held out and evaluated once: FEVER, DBpedia-entity, and the CQADupStack android and english forums. Two other forums, physics and programmers, select table penalties and routing settings and were reused during model construction.

Reference systems. bge-small-en-v1.5 and LEAF were selected for the project’s original evaluations. Each uses its own document representation. On BEIR-15 they score 0.5171 and 0.5402 against Nano’s 0.5081, at 2.57 and 1.39 ms median encoding; BM25 scores 0.4006. Nano’s broad-average score is below both dense references; the registered partition-specific contrasts follow.

Nano’s registered contrasts used a fixed sequence with one-sided 2.5% lower bounds. “Established” means the bound cleared zero; “unresolved” means it did not; “descriptive” rows were registered as reports, not tests. Intervals resample queries with the models fixed.

Contrast	Partition	nDCG@10 difference	Outcome
Nano – bge-small	clean-four	+0.0176; lower bound +0.0037	established
Nano – bge-small	all six	+0.0274; lower bound +0.0173	established
Nano – LEAF	all six	+0.0162; lower bound +0.0065	established
Nano – LEAF	clean-four	−0.0011; lower bound −0.0145	unresolved
Nano – bge-small	held-out NDO-3, dataset weighted	+0.0032 [−0.0069, +0.0134]	descriptive
Nano – bge-small	held-out NDO-3, query pooled	−0.0121 [−0.0219, −0.0024]	descriptive
Nano – LEAF	held-out NDO-3	−0.0389 [−0.0488, −0.0290]	descriptive

NDO-3 is the registered held-out aggregate: DBpedia-entity weighted one half, the android and english forums one quarter each, FEVER excluded as the double-exposure sensitivity row.

Zero’s registered contrasts used a Holm correction across the family. Zero missed its dense bar and did not confirmatorily beat BM25.

Contrast, all six	nDCG@10 difference	Raw 95% interval	Outcome
Zero – LightRetriever dense	−0.0243	[−0.0405, −0.0086]	below the 0.4583 bar
Zero – BM25	+0.0165	[+0.0017, +0.0311]	unresolved; p=.0149 fails Holm .0083
Zero + BM25 (convex) – OpenSearch	+0.0043	[−0.0063, +0.0151]	unresolved

The registered clean-four sensitivity is descriptive: −0.0443 [−0.0675, −0.0212], −0.0311 [−0.0517, −0.0109], and −0.0107 [−0.0262, +0.0043]. Unresolved superiority does not establish equivalence. The confirmatory fusion operator was convex score fusion with development-fitted weight 0.8 at depth 1000; the later DBSF-at-100 recommendation is a deployability choice, not a new confirmation. On the development components DBSF minus convex is +0.0035 at 10 candidates, −0.0020 at 50, −0.0063 at 100, and −0.0146 at 1000.

The held-out four, copied from their one-shot published aggregate:

System	FEVER	DBpedia	CQADup android	CQADup english
Nano	0.6231	0.4190	0.4774	0.4298
Zero	0.6978	0.3900	0.4431	0.3690
Stella	0.8207	0.4603	0.5275	0.5183

System	FEVER	DBpedia	CQADup android	CQADup english
bge-small	0.8662	0.4003	0.4761	0.4558
LEAF	0.8671	0.4520	0.5260	0.4707
BM25	0.5036	0.3045	0.3952	0.3481
Zero + BM25	0.7559	0.4030	0.4629	0.4016
Nano + BM25	0.6879	0.4296	0.4750	0.4323

Precision. All released paths run in fp32 through the published loader; a candidate fp16 export passed CPU parity and failed on CUDA (minimum cosine 0.662), which is why qualification runs through the actual loader and device.

Appendix B. Construction details

The roster. The 26 encoder spaces of Section 3.1, by width and own six-set score.

Encoder (frozen space)	Width	Pooling	Own nDCG@10	Table student	Head student	Roster
bge-small-en-v1.5	384	cls	0.5042	0.3581	0.5084	exploratory
arctic-embed-s	384	cls	0.4993	0.3516	0.2594	exploratory
gte-small	384	mean	0.4838	0.2903	0.4733	exploratory
arctic-embed-xs	384	cls	0.4662	0.3440	0.2779	exploratory
e5-small-v2	384	mean	0.4544	0.3203	0.3367	exploratory
minilm-l12	384	mean	0.4219	0.3138	0.3229	exploratory
minilm-l6	384	mean	0.4142	0.3267	0.3088	exploratory
multi-qa-minilm-l6	384	mean	0.4008	0.3406	0.3094	exploratory
msmarco-minilm-l6	384	mean	0.3281	0.2905	0.1962	exploratory
gte-base-en-v1.5	768	cls	0.5331	0.3252	0.2618	registered
arctic-embed-m-v1.5	768	cls	0.5263	0.3279	0.1874	registered
bge-base-en-v1.5	768	cls	0.5259	0.3529	0.3249	registered
bge-base-en-v1	768	cls	0.5131	0.3004	0.2773	exploratory
gte-base	768	mean	0.5063	0.2762	0.2314	exploratory
e5-base-v1	768	mean	0.4934	0.3126	0.2345	exploratory
e5-base-v2	768	mean	0.4669	0.2930	0.2241	registered
contriever-msmarco	768	mean	0.3909	0.3270	0.1879	exploratory
tas-b	768	cls	0.3262	0.2115	0.0868	exploratory
contriever	768	mean	0.2846	0.1864	0.0590	exploratory
gte-large-en-v1.5	1024	cls	0.5970	0.2455	0.2229	registered

Encoder (frozen space)	Width	Pooling	Own nDCG@10	Table student	Head student	Roster
stella-400M-v5	1024	mean	0.5745	0.3974	0.2505	registered
mxbai-embed-large-v1	1024	cls	0.5368	0.2605	0.2641	registered
bge-large-en-v1.5	1024	cls	0.5329	0.2845	0.2679	registered
arctic-embed-l	1024	cls	0.5289	0.3034	0.1564	registered
gte-large	1024	mean	0.5129	0.2481	0.2249	exploratory
e5-large-v2	1024	mean	0.4735	0.2585	0.1743	registered

Table student. Ridge regression from WordPiece token counts to the index’s query vectors, initialized from the index’s own token embeddings, solved by block conjugate gradient with a convergence gate; penalty grid 1e-4 to 1e-1 with one decade of extension at an edge, chosen on the two development forums and frozen before six-set scoring. The fit list was cleaned after an earlier version had 1.31% held-out-query overlap. One checkpoint failed the convergence gate and is excluded. Registered checkpoint exposure is documented; the exploratory expansion was not individually audited.

Head student. Frozen bge-small at the Nano build’s pinned revision; masked-mean-pooled hidden states of layers 12, 8, and 4 concatenated with a bias column; dense ridge with penalty scaled by the feature Gram trace, same grid, same selection order, same freeze; output normalized. bge-small-en-v1.5 is also an index in the roster and shares the backbone, and gte-small is nearly its linear image (retention 1.008 and 0.978); correlations without the backbone are +0.09 against the index and +0.46 against the table.

Family-adjusted quality model. Adding eight family indicators to the quality-and-width model raises in-sample R^2 to 0.91 (head) and 0.81 (table), preserving both coefficient signs. These models describe the observed families and are reported in sample. In Table B, the head’s standardized width and quality coefficients are -0.83 and $+0.69$; the table’s are -0.64 and $+0.63$. The table coefficient intervals overlap in magnitude.

Published retention context. LightRetriever reports about 95% retention for its jointly trained lookup path, while Zero retains 81.4% over frozen Stella here. These figures describe different models and evaluation protocols. Their comparison leaves the quality cost attributable to freezing the index unresolved.

Prospective roster (E24). Eight encoders fixed in the method before any encoding, all passing the tokenizer check: e5-large (1024), bge-large-en v1 (1024), e5-base-unsupervised (768), nomic-embed-text-v1 (768), msmarco-bert-base-dot-v5 (768), msmarco-distilbert-base-v4 (768), all-MiniLM-L12-v1 (384), paraphrase-MiniLM-L6-v2 (384). Teacher scores were measured first, the prediction file was written and its sha256 recorded (81f2c88e...) before any student was fitted, and every student command verified that hash. The table fit for e5-base-unsupervised stopped at the convergence gate at the extended penalty 1e-5; its head fit completed. E24 amendment 1 includes each completed fit, giving eight head points and seven table points. Per encoder, six-set score of the space, then table student actual / predicted, then head student actual / predicted:

Encoder	Width	Own nDCG@10	Table: actual / predicted	Head: actual / predicted
bge-large-en v1	1024	0.5237	0.2613 / 0.2883	0.2481 / 0.2103
e5-large	1024	0.4914	0.2875 / 0.2757	0.2243 / 0.1825
nomic-embed-text-v1	768	0.4851	0.3684 / 0.2945	0.2418 / 0.2336

Encoder	Width	Own nDCG@10	Table: actual / predicted	Head: actual / predicted
e5-base-unsupervised	768	0.4562	stopped / 0.2833	0.1178 / 0.2087
all-MiniLM-L12-v1	384	0.4115	0.3256 / 0.3171	0.3340 / 0.3062
msmarco-distilbert-base-v4	768	0.3436	0.3071 / 0.2395	0.1734 / 0.1117
msmarco-bert-base-dot-v5	768	0.3239	0.2127 / 0.2318	0.1032 / 0.0947
paraphrase-MiniLM-L6-v2	384	0.3001	0.2566 / 0.2738	0.2447 / 0.2103

msmarco-bert-base-dot-v5 is the teacher QED distilled from; its closed-form students retain 66% (table) and 32% (head) of its six-set score here, against the 92.5% QED reports for its trained two-layer student on BEIR under its own protocol, which the comparison does not reproduce.

Correlation rosters. Registered ten: Spearman head-versus-index +0.18 (−0.56 to +0.67), head-versus-table +0.24 (−0.49 to +0.74), head development-versus-public +0.71 (+0.12 to +0.96). Leave-one-family-out values range from +0.44 to +0.60 for head-versus-table and −0.12 to +0.27 for head-versus-index over the 26. Within width bands, head-versus-index Spearman is +0.47 [−0.47, +1.00] at 384 (n=9), +0.67 [−0.03, +1.00] at 768 (n=10), and +0.36 [−0.76, +0.87] at 1024 (n=7).

Vector fidelity and retrieval quality. Tables fitted to two e5 checkpoints match their index’s query vectors at mean cosine 0.90 and 0.89 against 0.78 for Stella’s table, yet retrieve worse. Three earlier training screens reduced their losses while leaving ranking unimproved; three vector-error diagnostics also failed to rank table quality. These observations motivate selection by retrieval against the target document vectors.

Costs. Table fitting four to seven minutes per registered index on an A100, including query encoding except cached Stella targets; head features 43 seconds for the fit list, dev grid 14 seconds per index, six-set scoring about one minute per index. Nano’s final optimization ran 57.3 hours on one A100 for 199,999,721 example presentations, about \$95 at the recorded rental rate, excluding target preparation, data generation, recipe search, failed runs, and evaluation. Zero’s retraining was estimated at 20 minutes with prepared targets, plus 8 to 12 hours to encode targets for a new index.

Appendix C. Search and precision controls

Instance sweep. Qdrant 1.19.1, HNSW m=16, ef_construct=100, indexing threshold 1 KB so no segment is served by a plain scan, ef in 16 to 512, oversampling 1, 2, and 4 under compression, exact parity against NumPy on 200 queries per path. Timing is warmed and sequential with encoding and search measured in separate phases. At ef=16 the uncompressed million-passage losses are 4.4%, 7.0%, and 16.3% for Stella, Nano, and Zero; under binary quantization without oversampling 6.1%, 9.2%, and 19.8%.

Cross-index sweep. The same engine and graph parameters, uncompressed, two builds per index and workload differing in a seeded insertion order, 16 to 512 ef, limit 11 with the self-hit drop, tie-aware exact parity on 200 queries per path. Four amendments were recorded during the run, before any index’s rows were analysed: SCIDOCs added as a third workload because TREC-COVID has 50 queries; the tie-aware measure; a 0.995 sanity gate on parity with all values reported; and the exclusion rule that removed tas-b (SCIDOCs head parity 0.994). Stella’s own-path rows agree with the instance sweep at every ef. Loss-based censoring: four indexes on FiQA, eight on SCIDOCs, eight on TREC-COVID; recovery-based censoring: one per workload.

Truncated queries (registered). On synthetic prefixes of SciFact, NFCorpus, and FiQA queries, the three dense paths retain descriptively similar shares of their own full-query nDCG@10 at each prefix length; this is not an equivalence test and no search-as-you-type session was measured.

Gap predictors (E22). Features per space and workload, declared before computation: the mean cosine distance from a query to its nearest document divided by the mean nearest-neighbor distance among a 5,000-document sample (for the student’s and the encoder’s queries, and their difference); the Gini coefficient of document occurrence counts across the student’s exact top-10 lists over the full document support, minus the encoder’s; the mean overlap of the student’s and encoder’s exact top-10; the effective rank (participation ratio) of a seeded 20,000-row sample of the encoder’s fit-query vectors and the ratio of the student’s to the encoder’s workload-query effective rank; the student-minus-encoder top-1 cosine and margin recorded by the search sweep; log2 width and log10 corpus size as covariates. Ridge with unit penalty on standardized features; the baseline is each training fold’s mean. The seven features not in Table F, same columns:

Feature	Spearman, 75 points	95% interval	FiQA / SCIDOCS / TREC-COVID	Leave-one-family- out Spearman, MAE (pp)
Student minus encoder ratio	−0.23	[−0.42, −0.03]	−0.37 / −0.44 / +0.17	−0.15, 2.23
Hubness of the student’s top-10, minus the encoder’s	+0.40	[+0.20, +0.57]	+0.12 / −0.20 / −0.15	+0.14, 2.11
Student-encoder top-10 overlap	+0.37	[+0.14, +0.60]	−0.01 / +0.43 / +0.28	+0.11, 2.29
Effective rank of the encoder’s fit-query cloud	−0.17	[−0.40, +0.11]	−0.28 / −0.34 / −0.05	−0.53, 2.38
Ratio of the student’s to the encoder’s workload-query effective rank	+0.18	[−0.06, +0.41]	+0.16 / −0.48 / +0.32	−0.30, 2.17
Top-1 cosine, student minus encoder	−0.05	[−0.30, +0.20]	0.00 / +0.23 / −0.34	−0.55, 2.28
Margin, student minus encoder	−0.28	[−0.48, −0.05]	−0.19 / −0.05 / −0.37	−0.11, 2.15

Corpus-size subsample (E23). FiQA documents drawn uniformly with seed 20260930 to 25,657, keeping every document with a positive judgment for a test query; each path’s exact top-10 and tie thresholds recomputed over the subsample; two builds; the same ef grid; per-space gaps paired against the full-corpus and SCIDOCS values with 10,000-draw bootstrap intervals over spaces.

Native binary pilot and replication. On one fixed binary collection per workload with 192 deterministically selected queries, binary scoring with rescoring at 1x oversampling loses 11.7, 5.1, and 2.3 percentage points of exact-neighbor recovery for Zero, Nano, and Stella on FiQA at ef=64; the million-passage replication loses 3.9, 3.1, and 2.2, and the extra Zero-versus-Nano interval includes zero. Native rescoring also changes cross-segment candidate merging.

Graph-free precision control. Same queries, vectors, and targets; all document sign codes scored exhaustively at global candidate budgets 10, 40, and 100; expected coverage averages uniform inclusion at tied boundaries.

Workload	Path	Sign query	Full query	Gain, points
FiQA	Zero	0.8217	0.9297	+10.80
FiQA	Nano	0.9137	0.9771	+6.33
FiQA	Stella	0.9638	0.9922	+2.84
1M diagnostic	Zero	0.9010	0.9677	+6.67
1M diagnostic	Nano	0.9484	0.9833	+3.49
1M diagnostic	Stella	0.9685	0.9932	+2.47

The paired extra Zero gains over Nano are 4.46 points on FiQA (95% query interval 2.53 to 6.43) and 3.17 on the diagnostic (1.58 to 4.75). Zero starts with more misses and does not repair a larger fraction of them. Full query values are not native scalar8 encoding, and global candidate counts are not per-segment oversampling.

Appendix D. Fusion and routing over the frozen index

Lexical fusion. Fusing Zero with BM25 raises its BEIR-15 macro score from 0.4572 to 0.4933, 87.9% of Stella; Nano plus BM25 reaches 0.5110. The sign varies by workload: Zero on HotpotQA rises from 0.6127 to 0.7015, Nano on MS MARCO falls from 0.4063 to 0.3699. On the million-passage diagnostic a dense-plus-sparse collection takes 4.71 ms per query including Zero encoding against 1.11 ms for the dense binary example. These timings compare different collection configurations. The scores use distribution-based score fusion over 100 candidates per branch with a bm25s lexical branch [DBSF]; other lexical implementations need their own comparison because their score distributions enter the normalization. The operator ordering changes with candidate depth (Appendix A).

Internal workload. Zero fell below BM25 on one internal domain-jargon workload. Three bounded attempts to specialize a student to that workload produced no improved encoder, as summarized in Section 7.

Routing and blending (registered). On 12 datasets, an oracle choosing Zero or Nano per query with relevance labels scores 0.5473 against 0.5161 for always-Nano, and matches always-Nano by sending the most beneficial 15% of queries to Nano. Query length, subwords per word, and the pooled vector norm are weak selectors (0.002 to 0.005 over random routing at matched Nano share); Zero’s post-search score margin recovers about a quarter of the oracle gain over random routing on six datasets and costs a second search for escalated queries. A development-selected 30% Zero contribution blended into Nano’s vector lowers Nano’s six-set score by 0.0051 (95% paired interval -0.0092 to -0.0011), consistent with how hard retrieval-strategy selection and query performance prediction remain [Arabzadeh et al., QPP].

References

- [QED] Yuxuan Wang and Hong Lyu. Query Encoder Distillation via Embedding Alignment is a Strong Baseline Method to Boost Dense Retriever Online Efficiency. SustainNLP Workshop at ACL, 2023. [arXiv:2306.11550](#).
- [EmbedDistill] Seungyeon Kim, Ankit Singh Rawat, Manzil Zaheer, et al. EmbedDistill: A Geometric Knowledge Distillation for Information Retrieval. 2023. [arXiv:2301.12005](#).
- [LEAF] Robin Vujanic and Thomas Rueckstiess. LEAF: Knowledge Distillation of Text Embedding Models with Teacher-Aligned Representations. 2025, revised 2026. [arXiv:2509.12539](#).
- [NanoVDR] Zhuchenyang Liu, Yao Zhang, and Yu Xiao. NanoVDR: Distilling a 2B Vision-Language Retriever into a 70M Text-Only Encoder for Visual Document Retrieval. 2026. [arXiv:2603.12824](#).
- [LightRetriever] Guangyuan Ma, Yongliang Ma, Xuanrui Gou, Zhenpeng Su, Ming Zhou, and Songlin Hu. LightRetriever: A LLM-based Text Retrieval Architecture with Extremely Faster Query Inference. ICLR, 2026. [arXiv:2505.12260](#).
- [pyNIFE] Stephan Tulkens. pyNIFE: Nearly Inference Free Embeddings in Python. Software, 2025. [Repository](#).
- [Model2Vec] MinishLab. Model2Vec: Fast State-of-the-Art Static Embeddings. Software. [Repository](#).

- [StaticEmb] Tom Aarsen. Train 400x faster Static Embedding Models with Sentence Transformers. Hugging Face, 15 January 2025. [Blog post](#).
- [Cho and Hariharan 2019] Jang Hyun Cho and Bharath Hariharan. On the Efficacy of Knowledge Distillation. ICCV, 2019. [Proceedings](#).
- [PROD] Zhenghao Lin, Yeyun Gong, Xiao Liu, et al. PROD: Progressive Distillation for Dense Retrieval. WWW, 2023. [arXiv:2209.13335](#).
- [HNSW] Yury A. Malkov and Dmitry A. Yashunin. Efficient and Robust Approximate Nearest Neighbor Search Using Hierarchical Navigable Small World Graphs. IEEE TPAMI 42(4), 2020. [arXiv:1603.09320](#).
- [OOD-DiskANN] Shikhar Jaiswal, Ravishankar Krishnaswamy, Ankit Garg, Harsha Vardhan Simhadri, and Sheshansh Agrawal. OOD-DiskANN: Efficient and Scalable Graph ANNS for Out-of-Distribution Queries. 2022. [arXiv:2211.12850](#).
- [RoarGraph] Meng Chen, Kai Zhang, Zhenying He, Yinan Jing, and X. Sean Wang. RoarGraph: A Projected Bipartite Graph for Efficient Cross-Modal Approximate Nearest Neighbor Search. PVLDB 17(11): 2735-2749, 2024. [Paper](#).
- [QA-Cos] Daehun Nyang. Beyond Hamming: Query-Aware Decoding of Binary Cosine Sketches. ICML 2026, PMLR 306: 94126-94143. [Proceedings](#).
- [Arabzadeh et al.] Negar Arabzadeh, Xinyi Yan, and Charles L. A. Clarke. Predicting Efficiency/Effectiveness Trade-offs for Dense vs. Sparse Retrieval Strategy Selection. CIKM, 2021. [arXiv:2109.10739](#).
- [QPP] Guglielmo Faggioli, Thibault Formal, Stefano Marchesin, Stéphane Clinchant, Nicola Ferro, and Benjamin Piwowarski. Query Performance Prediction for Neural IR: Are We There Yet? 2023. [arXiv:2302.09947](#).
- [DBSF] Qdrant contributors. Hybrid queries: distribution-based score fusion. Qdrant documentation. [Page](#).
- [Qdrant 1.19.1] Qdrant contributors. Binary query encoding and vector-index search implementation, v1.19.1. [Tagged source](#).
- [Stella] Stella-en-400M-v5. Model card for NovaSearch/stella_en_400M_v5. [Model page](#).
- [BEIR] BEIR contributors. BEIR retrieval benchmark, datasets, and evaluation software. [Repository](#).
- [Index migration] Qdrant contributors. Migrate to a New Embedding Model with Zero Downtime in Qdrant. Qdrant documentation. [Migration procedures](#).